

Bayesian Learning for Regression using Dirichlet Prior Distributions of Varying Localization

DISTRIBUTION A. Approved for public release: distribution unlimited.

Paul Rademacher¹ **Miloš Doroslovački²**

¹U.S. Naval Research Laboratory
Radar Division

²The George Washington University
Department of Electrical and Computer Engineering

July 11, 2021

Introduction

Bayesian approaches to statistical learning attempt to make better decisions by exploiting **prior knowledge** regarding the data-generating distribution:

Informative

- If the prior is localized around the true data-generating model, low-risk decisions can be made even with limited training data
- Priors that assign low weighting to the true model may not be able to realize satisfactory performance

Non-Informative

- Learners designed with minimally localized priors respond strongly to training data, avoiding the drawbacks of misinformed prior knowledge
- If the data volume is limited, high variance "overfit" solutions can occur

Dirichlet prior distributions have a number of desirable properties:

- **Full support** over the space of data-generating distributions, guaranteeing *consistent estimation* of the true data model
- They are **conjugate priors** for independent, identically distributed observations¹, leading to *closed-form* posterior distributions
- **Flexible parameterization** enabling *both* maximally and minimally informative priors and thus a wide range of learning solutions

¹Thomas S. Ferguson. "A Bayesian Analysis of Some Nonparametric Problems". In: *The Annals of Statistics* (1973).

Data Model and Regression Objective

Observable random variable: $x \in \mathcal{X} \subset \mathbb{R}$

Unobservable random variable: $y \in \mathcal{Y} \subset \mathbb{R}$

Observable training data: $D \in \mathcal{D} = \{\mathcal{Y} \times \mathcal{X}\}^N$

Independently, identically distributed according to an **unknown** probability mass function (PMF)

$$\theta \in \Theta = \left\{ \theta \in \mathbb{R}_{\geq 0}^{\mathcal{Y} \times \mathcal{X}} : \sum_{y \in \mathcal{Y}} \sum_{x \in \mathcal{X}} \theta(y, x) = 1 \right\},$$

such that $P_{y,x|\theta}(y, x|\theta) = P_{D_n|\theta}(y, x|\theta) = \theta(y, x)$.

Alternate Notation: $\theta \Leftrightarrow (\theta_m, \theta_c)$

- Marginal model $\theta_m \equiv \sum_{y \in \mathcal{Y}} \theta(y, \cdot) = P_{x|\theta} \equiv P_{x|\theta_m}$
- Conditional models $\theta_c(x) \equiv \theta(\cdot, x) / \theta_m(x) = P_{y|x,\theta} \equiv P_{y|x,\theta_c}$

Using the i.i.d. assumption,

$$P_{D|\theta}(D|\theta) = \left(\prod_{y \in \mathcal{Y}} \prod_{x \in \mathcal{X}} \theta(y, x)^{\Psi(y, x; D)} \right)^N$$

where data are represented using
 $\Psi : \mathcal{D} \mapsto \Psi \subset \Theta$, defined as

$$\Psi(y, x; D) = N^{-1} \sum_{n=1}^N \delta[(y, x), D_n] .$$

- Empirical distribution $\Psi(D)$ is a **sufficient statistic**² for the model θ
- Efficient: $|\Psi| = \binom{N+|\mathcal{Y}||\mathcal{X}|-1}{|\mathcal{Y}||\mathcal{X}|-1} \leq |\mathcal{D}|$

\Rightarrow **Represent data using new random process** $\psi \equiv \Psi(D) \in \Psi$

²Bernardo et al., *Bayesian Theory*.

- Conditioned on the true model, the data statistic is an "Empirical" random process
 $\psi|\theta \sim \text{Emp}(N, \theta)$
 - Equivalent to a normalized multinomial random process³
- As $N \rightarrow \infty$, the random process converges to $\psi|\theta \xrightarrow{p} \theta$
 \Rightarrow Use enables **consistent** estimation of model

Alternate Notation: $\psi \Leftrightarrow (\psi_m, \psi_c)$

- Marginal $\psi_m \equiv \sum_{y \in \mathcal{Y}} \psi(y, \cdot)$
- Conditional $\psi_c(x) \equiv \psi(\cdot, x) / \psi_m(x)$

By the aggregation property ⁴,

- $\psi_m | \theta_m \sim \text{Emp}(N, \theta_m)$
- $\psi_c(x) | \psi_m(x), \theta_c(x) \sim \text{Emp}(N \psi_m(x), \theta_c(x))$ are mutually **independent**

³Thomas P. Minka. *Bayesian inference, entropy, and the multinomial distribution*. Tech. rep. Microsoft Research, 2003.

⁴Johnson et al., *Discrete Multivariate Distributions*.

- Design a regression function $f : \Psi \mapsto \mathbb{R}^{\mathcal{X}}$ to minimize the expected squared-error with respect to θ :

$$\mathcal{R}_{\Theta}(f; \theta) = E_{y, x, \psi | \theta} \left[(f(x; \psi) - y)^2 \right] \equiv \underbrace{E_{x | \theta_m} [\Sigma_{y | x, \theta_c}]}_{\mathcal{R}_{\Theta}^*(\theta)} + \underbrace{E_{x, \psi | \theta} \left[(f(x; \psi) - \mu_{y | x, \theta_c})^2 \right]}_{\mathcal{R}_{\Theta, \text{ex}}(f; \theta)}$$

- Clairvoyant⁵ regressor $f_{\Theta}(x; \theta_c) = \mu_{y | x, \theta_c}$ achieves *irreducible* squared-error $\mathcal{R}_{\Theta}^*(\theta)$
- Excess squared-error can be decomposed into **bias** and **variance** terms:

$$\mathcal{R}_{\Theta, \text{ex}}(f; \theta) \equiv E_{x | \theta_m} \left[\left(E_{\psi | \theta} [f(x; \psi)] - f_{\Theta}(x; \theta_c) \right)^2 + C_{\psi | \theta} [f(x; \psi)] \right]$$

⁵Steven M. Kay. *Fundamentals of Statistical Signal Processing: Detection Theory*. Vol. 2. Prentice-Hall, 1998.

Model unknown. Select prior p_θ and formulate Bayesian risk:

$$\mathcal{R}(f) = E_\theta [\mathcal{R}_\Theta(f; \theta)] = E_{y, x, \psi} [(f(x; \psi) - y)^2]$$

$\Downarrow \quad \Downarrow \quad \Downarrow \quad \Downarrow$

Bayes optimal regressor:

$$f^*(x; \psi) = \arg \min_{y' \in \mathbb{R}} E_{y | x, \psi} [(y' - y)^2] = \mu_{y | x, \psi}$$

* Observe that $P_{y | x, \psi} = E_{\theta | x, \psi} [P_{y | x, \theta}] \equiv \mu_{\theta_c(x) | x, \psi}$

Bayesian distribution is the posterior mean⁶ of the predictive model θ_c

⁶Kevin P. Murphy. *Binomial and multinomial distributions*. Tech. rep. University of British Columbia, 2006.

Distributions: Prior to Predictive



The probability density function of the model $\theta \in \Theta$ is Dirichlet:

$$p_{\theta}(\theta) = \text{Dir}(\theta; \alpha_0, \alpha) \equiv \beta(\alpha_0 \alpha)^{-1} \prod_{y \in \mathcal{Y}} \prod_{x \in \mathcal{X}} \theta(y, x)^{\alpha_0 \alpha(y, x) - 1}$$

- Parameter α_0 controls localization around mean α

Alternate Notation:

- Marginal $\alpha_m \equiv \sum_{y \in \mathcal{Y}} \alpha(y, \cdot)$
- Conditional $\alpha_c(x) \equiv \alpha(\cdot, x) / \alpha_m(x)$
- * By the aggregation property⁷, $\theta_m \sim \text{Dir}(\alpha_0, \alpha_m)$ and $\theta_c(x) \sim \text{Dir}(\alpha_0 \alpha_m(x), \alpha_c(x))$ are mutually **independent**



⁷Ferguson, "A Bayesian Analysis of Some Nonparametric Problems".

Since $\perp\!\!\!\perp_x \theta_c(x)$ and $\theta_c \perp\!\!\!\perp \theta_m$, and since the Empirical process $\theta_c(x) | \psi_m(x), \psi_c(x)$ has exponential form, Dirichlet process $\theta_c(x)$ is conjugate⁸ and thus

$$\theta_c(x) | \psi_m(x), \psi_c(x) \sim \text{Dir} \left(\alpha_0 \alpha_m(x) + N \psi_m(x), \mu_{\theta_c(x) | \psi_m(x), \psi_c(x)} \right),$$

with mean functions


$$\mu_{\theta_c(x) | \psi_m(x), \psi_c(x)} = \gamma(x; \psi_m) \alpha_c(x) + (1 - \gamma(x; \psi_m)) \psi_c(x) \equiv \mathbf{P}_{y|x,\psi},$$

where $\gamma(x; \psi_m) = \left(1 + N \psi_m(x) / (\alpha_0 \alpha_m(x)) \right)^{-1} \in (0, 1]$.

Bayesian predictions mix prior mean α_c with empirical distribution ψ_c

⁸Sergios Theodoridis. *Machine Learning: A Bayesian and Optimization Perspective*. Elsevier, 2015.

- As localization α_0 increases,
 $\theta_c(x) | \psi_m(x), \psi_c(x) \xrightarrow{p} \alpha_c(x)$ and the prior is emphasized
- As training volume N increases,
 $\theta_c(x) | \psi_m(x), \psi_c(x) \xrightarrow{p} \psi_c(x)$ and data is emphasized
 - Since $\psi_c | \theta_c \xrightarrow{p} \theta_c$, the true predictive model is **identified**



Full support prior ensures consistent estimation of model

Regressor: $f^*(x; \psi) \equiv \gamma(x; \psi_m) \mu_{y|x} + (1 - \gamma(x; \psi_m)) \sum_{y \in \mathcal{Y}} \psi_c(y; x) y$

- * Convexly combines first moment of $P_{y|x} = \mu_{\theta_c(x)} = \alpha_c(x)$ with empirical mean
- * Inherits trends from posterior distribution, allowing maximal or minimal confidence in the prior

Excess SE: $\mathcal{R}_{\Theta, \text{ex}}(f^*; \theta) \equiv E_{x|\theta_m} \left[\lambda_{\text{Bias}}(x; \theta_m) (\mu_{y|x} - \mu_{y|x, \theta_c})^2 + \lambda_{\text{Var}}(x; \theta_m) \Sigma_{y|x, \theta_c} \right]$

- * $\lambda_{\text{Bias}}(x; \theta_m) = E_{\psi_m | \theta_m} [\gamma(x; \psi_m)^2]$ and $\lambda_{\text{Var}}(x; \theta_m) = E_{\psi_m | \theta_m} \left[\frac{(1 - \gamma(x; \psi_m))^2}{N \psi_m(x)} \right]$
- **Bias:** proportionate to squared-difference between data-independent regressor $\mu_{y|x}$ and clairvoyant regressor
- **Variance:** proportionate to the predictive variance, adding to the irreducible risk

Trends and Results

Data Model:



- $\mathcal{X} = \mathcal{Y} = \{i/127 : i = 0, \dots, 127\}$
- $\theta_m = |\mathcal{X}|^{-1}$
- **Clairvoyant:** $\mu_{y|x, \theta_c} = 1/(2 + \sin(2\pi x))$
- $\Sigma_{y|x, \theta} = 0.2 \mu_{y|x, \theta} (1 - \mu_{y|x, \theta})$
 $\Rightarrow \mathcal{R}_{\Theta}^*(\theta) \approx 0.039$

Learners:

- Dirichlet:
 - $\alpha_m = |\mathcal{X}|^{-1}$
 - **Prior** $\mu_{y|x} = 0.5$
- Normal⁹:
 - $y|x, \theta \sim \mathcal{N}([1, x]\theta, 0.1)$
 - $\theta \sim \mathcal{N}([0.5, 0], \Sigma_{\theta})$

- *Prior confidence* of Dirichlet and Normal learners varied using α_0 and Σ_{θ}
- Both learners effect the same **biased** untrained regressor to approximate the non-linear clairvoyant regressor

⁹Theodoridis, *Machine Learning: A Bayesian and Optimization Perspective*.



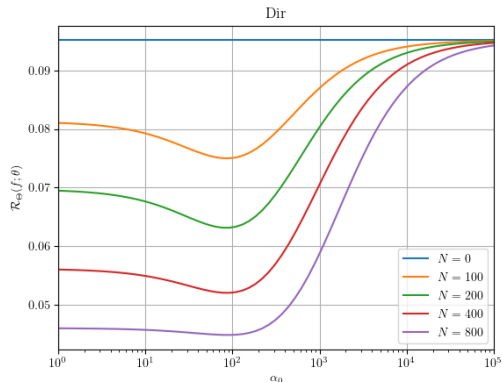
- Lines show bias $E_{\psi|\theta} [\mu_{y|x,\psi}]$, fill regions shows variance $C_{\psi|\theta} [\mu_{y|x,\psi}]$
- Python simulation results average 50,000 learning iterations

For a given conditional mean α_c ,
localization $\bar{\alpha}_0(x) \equiv \alpha_0 \alpha_m(x)$ controls a
Bias-Variance trade-off:


$\bar{\alpha}_0(x)$	$\lambda_{\text{Bias}}(x; \theta_m)$	$\lambda_{\text{Var}}(x; \theta_m)$
$\rightarrow \infty$	1	0
$\rightarrow 0$	$(1 - \theta_m(x))^N$	$E_{\Psi_m \theta_m} \left[(N \Psi_m(x))^{-1} \right]$

Optimal:

$$\bar{\alpha}_0(x) = \frac{\Sigma_{y|x, \theta_c}}{(\mu_{y|x} - \mu_{y|x, \theta_c})^2}$$



- As $N \rightarrow \infty$, both $\lambda_{\text{Bias}}(\mathbf{x}; \theta_m) \rightarrow 0$ and $\lambda_{\text{Var}}(\mathbf{x}; \theta_m) \rightarrow 0$
 $\Rightarrow \mathcal{R}_{\Theta, \text{ex}}(f^*; \theta) \rightarrow 0$ for **any** model θ
- Note that $f^*(\mathbf{x}; \psi)$ converges to the clairvoyant regressor **regardless** of how biased the prior conditional mean α_c is, or how much confidence in α_c is indicated through the localization α_0



Summary

Full-support Bayesian learning with a Dirichlet prior enables:

- Asymptotically **optimal** performance for data-rich applications
- **Maximal** prior knowledge required for data-limited applications

Future Work

- Generalize these concepts for more general data models using the continuous Dirichlet process¹⁰
 - Practical necessity motivates the use of **discretization** to realize the demonstrated benefits
- Use the Dirichlet prior with different likelihood functions (e.g., mixture model) to effect limited-support priors that may be best suited for data-limited applications

¹⁰Samuel J. Gershman et al. "A tutorial on Bayesian nonparametric models". In: *Journal of Mathematical Psychology* 56 (2012).